

```
In [4]: import pandas as pd
df = pd.read_excel("C:/Users/rohit/Downloads/transactions-2025-07-23.csv.xlsx")
df.head()
```

```
Out[4]:
```

	TRANSACTION_NUMBER	INSTANCE_DATE	GROUP_EN	PROCEDURE_EN	IS_OFFPLAN_EN
0	102-1-2025	2025-01-07 15:57:40	Sales	Sell - Pre registration	Off-Plan
1	102-10-2025	2025-01-02 08:23:00	Sales	Sell - Pre registration	Off-Plan
2	102-100-2025	2025-01-02 15:18:51	Sales	Sell - Pre registration	Off-Plan
3	102-1000-2025	2025-01-08 11:02:57	Sales	Sell - Pre registration	Off-Plan
4	102-10000-2025	2025-02-07 11:15:01	Sales	Sell - Pre registration	Off-Plan

5 rows × 22 columns



```
In [5]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 110490 entries, 0 to 110489
Data columns (total 22 columns):
#   Column                Non-Null Count  Dtype
---  -
0   TRANSACTION_NUMBER    110490 non-null object
1   INSTANCE_DATE         110490 non-null datetime64[ns]
2   GROUP_EN              110490 non-null object
3   PROCEDURE_EN          110490 non-null object
4   IS_OFFPLAN_EN         110490 non-null object
5   IS_FREE_HOLD_EN       110490 non-null object
6   USAGE_EN              110490 non-null object
7   AREA_EN               110490 non-null object
8   PROP_TYPE_EN          110490 non-null object
9   PROP_SB_TYPE_EN       104394 non-null object
10  TRANS_VALUE           110490 non-null float64
11  PROCEDURE_AREA        110489 non-null float64
12  ACTUAL_AREA           110490 non-null float64
13  ROOMS_EN              94277 non-null object
14  PARKING               85646 non-null object
15  NEAREST_METRO_EN      56461 non-null object
16  NEAREST_MALL_EN       55007 non-null object
17  NEAREST_LANDMARK_EN   68028 non-null object
18  TOTAL_BUYER           110490 non-null int64
19  TOTAL_SELLER          110490 non-null int64
20  MASTER_PROJECT_EN     277 non-null object
21  PROJECT_EN            100900 non-null object
dtypes: datetime64[ns](1), float64(3), int64(2), object(16)
memory usage: 18.5+ MB
```

```
In [6]: #Columns with more than 30% null values
drop_cols = ['MASTER_PROJECT_EN', 'NEAREST_METRO_EN', 'NEAREST_MALL_EN', 'NEAREST_L
df_cleaned = df.drop(columns=drop_cols)
```

```
In [7]: #Null audit - find percentage of missing values
null_summary = df.isnull().mean().sort_values(ascending=False)
print(null_summary[null_summary > 0])
```

```
MASTER_PROJECT_EN    0.997493
NEAREST_MALL_EN       0.502154
NEAREST_METRO_EN      0.488994
NEAREST_LANDMARK_EN   0.384306
PARKING               0.224853
ROOMS_EN              0.146737
PROJECT_EN            0.086795
PROP_SB_TYPE_EN       0.055172
PROCEDURE_AREA        0.000009
dtype: float64
```

```
In [8]: #impute missing values for numerical columns
median_val = df_cleaned['PROCEDURE_AREA'].median()
df_cleaned.loc[:, 'PROCEDURE_AREA'] = df_cleaned['PROCEDURE_AREA'].fillna(median_val)
```

```
In [9]: #impute missing values for categorical columns
categorical_impute = ['PARKING', 'ROOMS_EN', 'PROJECT_EN', 'PROP_SB_TYPE_EN']
for col in categorical_impute:
```

```
mode_val = df_cleaned[col].mode()[0]
df_cleaned.loc[:, col] = df_cleaned[col].fillna(mode_val)
```

```
In [10]: #check missing values
print(df_cleaned.isnull().sum().sort_values(ascending=False).head())
```

```
TRANSACTION_NUMBER    0
INSTANCE_DATE          0
GROUP_EN              0
PROCEDURE_EN          0
IS_OFFPLAN_EN         0
dtype: int64
```

```
In [11]: import numpy as np
import matplotlib.pyplot as plt
from matplotlib.ticker import FuncFormatter, LogLocator

# -----
# 1) Data prep & key stats
# -----
s = df_cleaned['TRANS_VALUE'].dropna()
s = s[s > 0] # Ensure compatibility with log scale

median_val = s.median()
p99_val     = s.quantile(0.99)

bins = np.logspace(np.log10(s.min()), np.log10(s.max()), 80)

# -----
# 2) Ivy-League Styling
# -----
plt.style.use('seaborn-v0_8-white')
plt.rcParams.update({
    "font.family": "serif",
    "axes.labelsize": 12,
    "axes.titlesize": 15,
    "figure.dpi": 120,
    "axes.edgecolor": "#2F2F2F",
    "axes.linewidth": 1.1
})

def aed_millions(x, pos):
    return f"AED {x/1e6:.0f}M" if x >= 1e6 else f"AED {x:,.0f}"

# -----
# 3) Plot
# -----
fig, ax = plt.subplots(figsize=(11, 6))

# Main histogram with a muted navy tone
ax.hist(s, bins=bins, histtype='stepfilled',
        color='#1A2A5E', alpha=0.8, edgecolor='black')

# Highlight top 1% tail in soft gold
ax.axvspan(p99_val, s.max(), color='#C49E47', alpha=0.15, label='Top 1% tail')
```

```

# Vertical reference lines
ax.axvline(median_val, color='#B03060', linestyle='--', linewidth=2,
           label=f"Median ≈ AED {median_val:,.0f}")
ax.axvline(p99_val, color='#C49E47', linestyle='--', linewidth=2,
           label=f"99th pct ≈ AED {p99_val:,.0f}")

# Log x-axis with AED formatting
ax.set_xscale('log')
ax.xaxis.set_major_locator(LogLocator(base=10.0, numticks=10))
ax.xaxis.set_major_formatter(FuncFormatter(aed_millions))


# Titles & Labels
ax.set_title("Distribution of Transaction Values (AED) – Log Scale", fontsize=16, f
ax.set_xlabel("Transaction Value (AED, log scale)")
ax.set_ylabel("Frequency")

# Legend & grid
ax.legend(frameon=True, fontsize=10)
ax.grid(which="both", linestyle=":", linewidth=0.6, alpha=0.5)

plt.tight_layout()
plt.show()

print(f"Median Transaction Value: AED {median_val:,.0f}")
print(f"99th Percentile Transaction Value: AED {p99_val:,.0f}")

```



Median Transaction Value: AED 1,590,000

99th Percentile Transaction Value: AED 21,000,000

```

In [12]: import numpy as np
import matplotlib.pyplot as plt
from matplotlib.ticker import FuncFormatter, LogLocator

# -----
# 1) Data Preparation
# -----
area = df_cleaned['ACTUAL_AREA'].dropna()

```

```

area = area[area > 0] # Log scale can't handle zero or negatives

median_area = area.median()
p99_area    = area.quantile(0.99)

bins = np.logspace(np.log10(area.min()), np.log10(area.max()), 80)

# -----
# 2) Professional Styling
# -----
plt.style.use('seaborn-v0_8-white')
plt.rcParams.update({
    "font.family": "serif",
    "axes.labelsize": 12,
    "axes.titlesize": 15,
    "figure.dpi": 120,
    "axes.edgecolor": "#2F2F2F",
    "axes.linewidth": 1.1
})

def sqm_formatter(x, pos):
    return f"{x:,.0f} m²"

# -----
# 3) Plot
# -----
fig, ax = plt.subplots(figsize=(11, 6))

# Main histogram – deep teal
ax.hist(area, bins=bins, histtype='stepfilled',
        color='#1B4D3E', alpha=0.85, edgecolor='black')

# Highlight top 1% tail – champagne gold
ax.axvspan(p99_area, area.max(), color='#D4AF37', alpha=0.2, label='Top 1% tail')

# Reference Lines
ax.axvline(median_area, color='#800020', linestyle='--', linewidth=2,
          label=f"Median ≈ {median_area:,.1f} m²")
ax.axvline(p99_area, color='#D4AF37', linestyle='--', linewidth=2,
          label=f"99th pct ≈ {p99_area:,.0f} m²")

# Log x-axis
ax.set_xscale('log')
ax.xaxis.set_major_locator(LogLocator(base=10.0, numticks=10))
ax.xaxis.set_major_formatter(FuncFormatter(sqm_formatter))

# Titles and Labels
ax.set_title("Distribution of Property Areas (m²) – Log Scale", fontsize=16, fontwe
ax.set_xlabel("Property Area (m², log scale)")
ax.set_ylabel("Frequency")


# Legend & grid
ax.legend(frameon=True, fontsize=10)
ax.grid(which="both", linestyle=":", linewidth=0.6, alpha=0.5)

plt.tight_layout()

```

```
plt.show()

# Key Stats
print(f"Median Property Area: {median_area:,.1f} m²")
print(f"99th Percentile Property Area: {p99_area:,.0f} m²")
```



Median Property Area: 86.9 m²

99th Percentile Property Area: 1,154 m²

```
In [13]: import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt

# Log transform transaction values (if not already done)
df_cleaned['LOG_TRANS_VALUE'] = np.log10(df_cleaned['TRANS_VALUE'].clip(lower=1))

prop_order = ['Unit', 'Building', 'Land'] # Adjust as needed
palette = {
    'Unit': '#2E86C1',      # Blue
    'Building': '#E67E22',  # Orange
    'Land': '#239B56'       # Green
}

plt.figure(figsize=(10, 6))
ax = sns.boxplot(
    data=df_cleaned,
    x='PROP_TYPE_EN',
    y='LOG_TRANS_VALUE',
    hue='PROP_TYPE_EN',
    order=prop_order,
    hue_order=prop_order,
    dodge=False,
    palette=palette
)

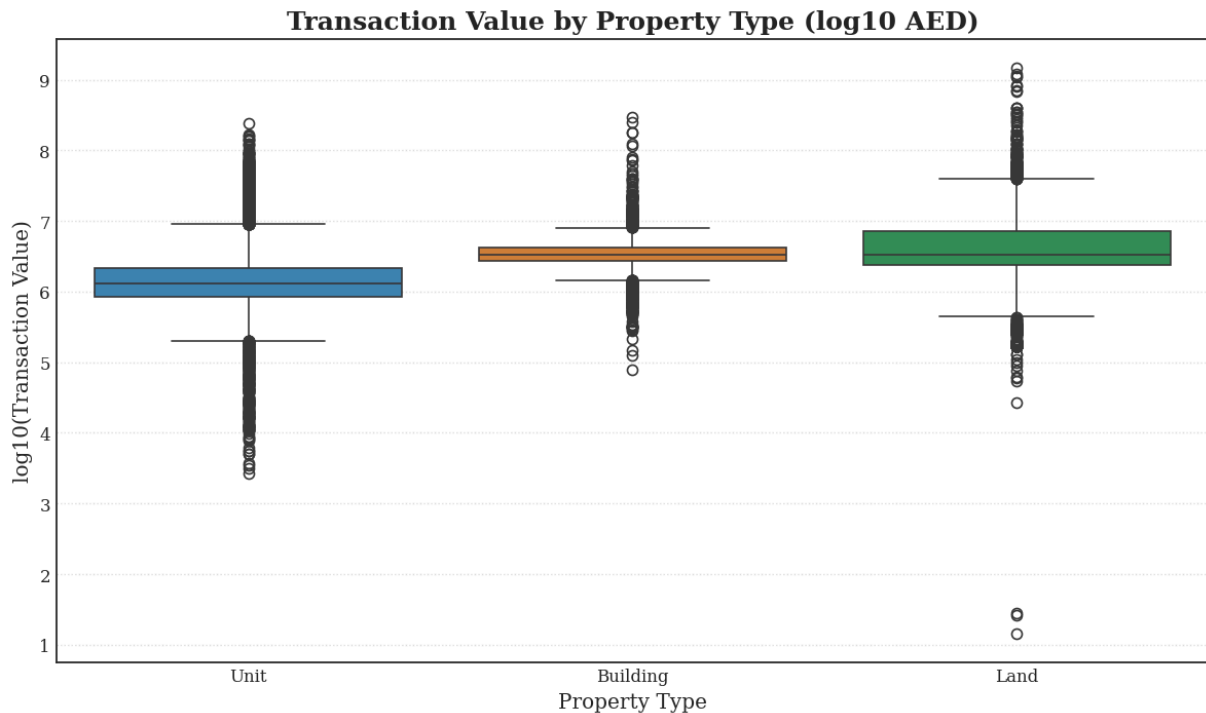
# Remove legend safely if it exists
if ax.get_legend() is not None:
```

```

ax.get_legend().remove()

plt.title("Transaction Value by Property Type (log10 AED)", fontsize=15, fontweight
plt.xlabel("Property Type")
plt.ylabel("log10(Transaction Value)")
plt.grid(axis='y', linestyle=':', alpha=0.6)
plt.tight_layout()
plt.show()

```



```

In [14]: import re
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt

# -----
# Prepare data
# -----
df_plot = df_cleaned.copy()

def to_numeric_bedrooms(x):
    if pd.isna(x): return np.nan
    s = str(x).lower().strip()
    if 'studio' in s: return 0
    digits = re.sub(r'^0-9', '', s)
    return float(digits) if digits else np.nan

df_plot['ROOMS_EN'] = df_plot['ROOMS_EN'].apply(to_numeric_bedrooms)
df_plot = df_plot.dropna(subset=['ROOMS_EN', 'TRANS_VALUE'])
df_plot = df_plot[df_plot['TRANS_VALUE'] > 0]
df_plot = df_plot[df_plot['ROOMS_EN'].between(0, 7)]

medians = df_plot.groupby('ROOMS_EN')['TRANS_VALUE'].median()
bedroom_order = sorted(df_plot['ROOMS_EN'].unique())

```

```

# -----
# Plot
# -----
plt.figure(figsize=(12, 6))
palette = sns.color_palette("coolwarm", len(bedroom_order))


ax = sns.boxplot(
    data=df_plot,
    x='ROOMS_EN',
    y='TRANS_VALUE',
    order=bedroom_order,
    hue='ROOMS_EN',
    dodge=False,
    legend=False,
    palette=palette
)

ax.set_yscale('log')

# Place median labels closer to the box, not on top of chart
for i, room in enumerate(bedroom_order):
    m = medians.loc[room]
    ax.text(
        i, m * 1.3, # Adjust multiplier to control height
        f"{m/1e6:.1f}M",
        ha='center', va='bottom',
        fontsize=10, fontweight='bold',
        bbox=dict(facecolor='white', edgecolor='none', alpha=0.8)
    )

ax.set_title("Transaction Value Distribution by Number of Bedrooms", fontsize=14, f
ax.set_xlabel("Number of Bedrooms")
ax.set_ylabel("Transaction Value [AED] (log scale)")
plt.tight_layout()
plt.show()

```




```
In [15]: import seaborn as sns
import matplotlib.pyplot as plt

# Select key numeric columns
df_corr = df_cleaned[['TRANS_VALUE', 'ACTUAL_AREA', 'ROOMS_EN']].copy()

# Convert ROOMS_EN to numeric if needed
def to_numeric_bedrooms(x):
    if pd.isna(x): return np.nan
    s = str(x).lower().strip()
    if 'studio' in s: return 0
    digits = ''.join([c for c in s if c.isdigit()])
    return float(digits) if digits else np.nan

df_corr['ROOMS_EN'] = df_corr['ROOMS_EN'].apply(to_numeric_bedrooms)

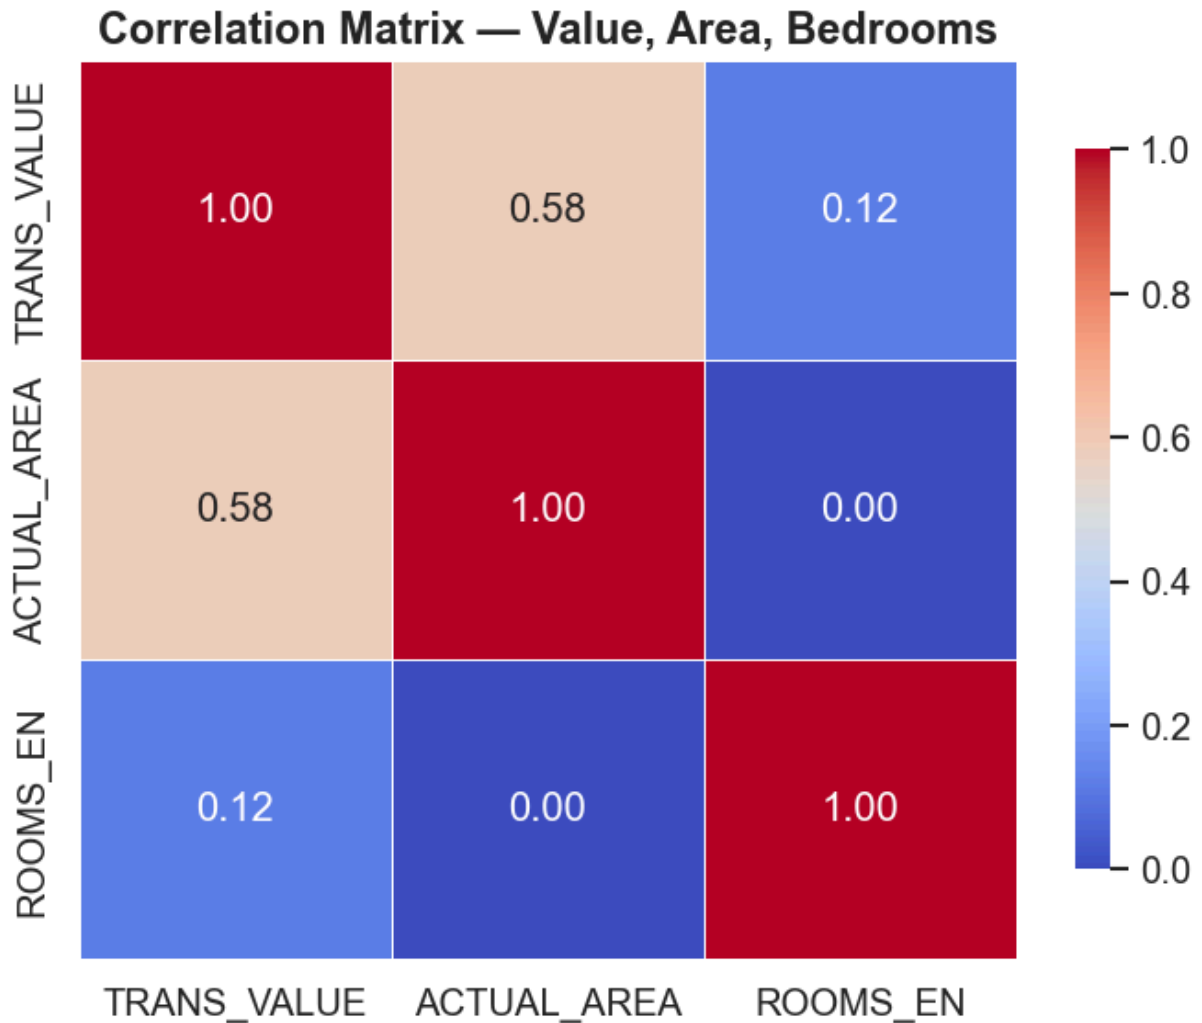
# Drop rows with missing values
df_corr = df_corr.dropna()

# Compute correlation matrix
corr_matrix = df_corr.corr()

# Plot heatmap
plt.figure(figsize=(6, 5))
sns.set_theme(style="white", font_scale=1.1)

ax = sns.heatmap(
    corr_matrix,
    annot=True,
    fmt=".2f",
    cmap="coolwarm",
    vmin=0, vmax=1,
    linewidths=0.5,
    linecolor='white',
    cbar_kws={"shrink": 0.8}
)

plt.title("Correlation Matrix – Value, Area, Bedrooms", fontsize=14, fontweight='bold')
plt.tight_layout()
plt.show()
```



```
In [16]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

# --- Monthly Data Preparation (same as before) ---
df_trends = df_cleaned.copy()
df_trends['INSTANCE_DATE'] = pd.to_datetime(df_trends['INSTANCE_DATE'])
df_2025 = df_trends[df_trends['INSTANCE_DATE'].dt.year == 2025].copy()
df_2025['Month'] = df_2025['INSTANCE_DATE'].dt.to_period('M').astype(str)

monthly_summary = df_2025.groupby('Month').agg(
    transactions=('TRANSACTION_NUMBER', 'count'),
    total_value=('TRANS_VALUE', 'sum')
).reset_index()
monthly_summary['total_value_billion'] = monthly_summary['total_value'] / 1e9

# Peak values
peak_idx = monthly_summary['transactions'].idxmax()
peak_month = monthly_summary.loc[peak_idx, 'Month']
peak_trans = monthly_summary.loc[peak_idx, 'transactions']
peak_value = monthly_summary.loc[peak_idx, 'total_value_billion']

# --- Plot Settings ---
sns.set_theme(style="whitegrid")
```

```

plt.rcParams.update({
    "font.family": "serif",
    "axes.titlesize": 15,
    "axes.labelsize": 13,
    "xtick.labelsize": 11,
    "ytick.labelsize": 11,
    "figure.dpi": 120,
})

# --- Larger Figure with Adjusted Layout ---
fig, axes = plt.subplots(1, 2, figsize=(16, 5))


# --- Monthly Transaction Volume ---
sns.lineplot(
    data=monthly_summary,
    x='Month',
    y='transactions',
    marker='o',
    color='navy',
    linewidth=2,
    ax=axes[0]
)
axes[0].set_title("Monthly Transaction Volume", fontsize=15, fontweight='bold', loc='right')
axes[0].set_xlabel("Month")
axes[0].set_ylabel("Number of Transactions")
axes[0].tick_params(axis='x', rotation=45)
axes[0].annotate(
    f"Peak: {peak_trans:,} sales",
    xy=(peak_idx, peak_trans),
    xytext=(peak_idx, peak_trans + 1200),
    fontsize=11, fontweight='bold',
    ha='center',
    arrowprops=dict(arrowstyle="->", color='black', linewidth=1.2)
)

# --- Monthly Total Transaction Value ---
sns.lineplot(
    data=monthly_summary,
    x='Month',
    y='total_value_billion',
    marker='o',
    color='darkgreen',
    linewidth=2,
    ax=axes[1]
)
axes[1].set_title("Monthly Total Transaction Value", fontsize=15, fontweight='bold', loc='right')
axes[1].set_xlabel("Month")
axes[1].set_ylabel("Total Value (AED, billions)")
axes[1].tick_params(axis='x', rotation=45)
axes[1].annotate(
    f"Peak in {peak_month}: {peak_value:.1f}B AED",
    xy=(peak_idx, peak_value),
    xytext=(peak_idx, peak_value + 6),
    fontsize=11, fontweight='bold',
    ha='center',
    arrowprops=dict(arrowstyle="->", color='black', linewidth=1.2)
)

```

```
)

# --- Reduce Top Gap ---
plt.suptitle("Dubai Real Estate — Monthly Trends (2025)", fontsize=17, fontweight='bold')
plt.tight_layout(rect=[0, 0, 1, 0.92]) # Less top padding
plt.show()
```



```
In [17]: import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt

# -----
# 1) Filter Jan-Jul 2025
# -----
df_geo = df_cleaned.copy()
df_geo['INSTANCE_DATE'] = pd.to_datetime(df_geo['INSTANCE_DATE'])

start, end = '2025-01-01', '2025-07-31'
mask = (df_geo['INSTANCE_DATE'] >= start) & (df_geo['INSTANCE_DATE'] <= end)
df_geo = df_geo.loc[mask]

# -----
# 2) Aggregate by AREA_EN
# -----
g = (df_geo
     .groupby('AREA_EN', dropna=False)
     .agg(
         transactions=('TRANSACTION_NUMBER', 'count'),
         total_value=('TRANS_VALUE', 'sum')
     )
     .sort_values('transactions', ascending=False))

# Calculate shares
total_txn = g['transactions'].sum()
g['txn_share_%'] = 100 * g['transactions'] / total_txn

total_val = g['total_value'].sum()
g['val_share_%'] = 100 * g['total_value'] / total_val

# -----
# 3) Top 10 by transactions
# -----
top10_txn = g.nlargest(10, 'transactions').reset_index()
```

```

print("Top 10 areas by number of transactions (Jan-Jul 2025)")
print(top10_txn[['AREA_EN', 'transactions', 'txn_share_%']].round({'txn_share_%': 2

# -----
# 4) Plot - Top 10 by transactions
# -----
sns.set_theme(style="whitegrid")
plt.figure(figsize=(10, 6))

# FIX: Use 'color' instead of 'palette' (no hue required)
ax = sns.barplot(
    data=top10_txn,
    y='AREA_EN',
    x='transactions',
    color='steelblue' # Fixed to avoid FutureWarning
)


# Annotate bars with counts & share
for i, r in top10_txn.iterrows():
    ax.text(r['transactions'] * 1.01, i,
           f"{r['transactions']:,} ({r['txn_share_%']:.1f}%)",
           va='center', fontsize=10)

plt.title("Top 10 Areas by Number of Transactions (Jan-Jul 2025)",
         fontsize=14, fontweight='bold')
plt.xlabel("Number of Transactions")
plt.ylabel("Area")
plt.tight_layout()
plt.show()

```

Top 10 areas by number of transactions (Jan-Jul 2025)

	AREA_EN	transactions	txn_share_%
0	JUMEIRAH VILLAGE CIRCLE	9873	8.94
1	BUSINESS BAY	6664	6.03
2	Al Yelayiss 1	5152	4.66
3	DUBAI LAND RESIDENCE COMPLEX	4117	3.73
4	Madinat Al Mataar	3691	3.34
5	MOTOR CITY	3168	2.87
6	DUBAI PRODUCTION CITY	3155	2.86
7	DUBAI MARINA	3000	2.72
8	JUMEIRAH LAKES TOWERS	2587	2.34
9	Bukadra	2402	2.17



```
In [18]: import matplotlib.ticker as ticker

# Top 10 by transaction value
top10_val = g.nlargest(10, 'total_value').reset_index()


# Plot 2: Transaction Value
plt.figure(figsize=(11, 7))
ax = sns.barplot(
    data=top10_val,
    y='AREA_EN',
    x='total_value',
    hue='AREA_EN',
    legend=False,
    palette='Reds_r'
)

def billions(x, pos):
    return f'{x/1e9:.0f}B'
ax.xaxis.set_major_formatter(ticker.FuncFormatter(billions))

# Annotate each bar
for i, r in top10_val.iterrows():
    ax.text(
        r['total_value'] * 1.01,
        i,
        f"AED {r['total_value']/1e9:.1f}B ({r['val_share_%']:.1f}%)",
        va='center',
        fontsize=10
    )

ax.set_title("Top 10 Areas by Total Transaction Value (H1 2025)", fontsize=14, font
ax.set_xlabel("Total Value (AED, Billions)")
ax.set_ylabel("Area")
```

```
plt.tight_layout()
plt.show()
```



```
In [19]: # 0) Imports

import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
import numpy as np # Used for checking numerical types
import sklearn # Import sklearn to check version if needed

# 1) Configuration
# - Centralize parameters for easy modification

FILE_PATH = "C:/Users/rohit/Downloads/transactions-2025-07-23.csv.xlsx"
TARGET_COLUMN = 'TRANS_VALUE'

# Columns to potentially drop if they have too many missing values
COLUMNS_TO_CHECK_MISSING = [
    'MASTER_PROJECT_EN', 'NEAREST_METRO_EN', 'NEAREST_MALL_EN',
    'NEAREST_LANDMARK_EN'
]
MISSING_THRESHOLD_PERCENT = 30 # Drop columns with over 30% missing values

# Specific columns for numerical imputation
NUMERICAL_IMPUTE_COLUMNS = ['PROCEDURE_AREA'] # Example if you have others

# Specific columns for categorical imputation
CATEGORICAL_IMPUTE_COLUMNS = ['PARKING', 'ROOMS_EN', 'PROJECT_EN', 'PROP_SB_TYPE_EN']
```

```

# Features to be used in the model
# Ensure these align with your imputation and data types
MODEL_FEATURES = ['PROCEDURE_AREA', 'PARKING', 'ROOMS_EN', 'PROJECT_EN', 'PROP_SB_T

RANDOM_STATE = 42
TEST_SIZE = 0.2
N_ESTIMATORS = 100

# 2) Load & Initial Data Inspection

print("Attempting to load data...")
try:
    # Check for the correct file type and use appropriate reader
    if FILE_PATH.lower().endswith('.csv.xlsx') or FILE_PATH.lower().endswith('.xlsx f
        df = pd.read_excel(FILE_PATH)
    elif FILE_PATH.lower().endswith('.csv'):
        df = pd.read_csv(FILE_PATH)
    else:
        raise ValueError("Unsupported file format. Please provide a .csv or .xlsx f
    print("✅ Data loaded successfully!")
    print(f"Initial DataFrame shape: {df.shape}")
    print("Initial DataFrame head:\n", df.head())
    print("Initial DataFrame info:")
    df.info() # Provides non-null counts and dtypes
    print("-" * 50)

except FileNotFoundError:
    print(f"❌ Error: File not found at '{FILE_PATH}'. Please check the path and f
    exit()
except Exception as e:
    print(f"❌ An error occurred during data loading: {e}")
    exit()

# Make a copy to avoid SettingWithCopyWarning
df_cleaned = df.copy()

# 3) Data Cleaning

print("\nStarting data cleaning...")

# Drop columns with over MISSING_THRESHOLD_PERCENT missing values
print(f"Checking columns to drop with > {MISSING_THRESHOLD_PERCENT}% missing values
cols_to_drop_actual = []
for col in COLUMNS_TO_CHECK_MISSING:
    if col in df_cleaned.columns:
        missing_percent = df_cleaned[col].isnull().sum() / len(df_cleaned) * 100
        if missing_percent > MISSING_THRESHOLD_PERCENT:
            cols_to_drop_actual.append(col)
            print(f" - Dropping '{col}' (missing: {missing_percent:.2f}%)")
    else:
        print(f" - Warning: Column '{col}' not found in DataFrame for missing chec

if cols_to_drop_actual:

```



```

df_cleaned = df_cleaned.drop(columns=cols_to_drop_actual)
print(f"✅ Dropped {len(cols_to_drop_actual)} columns due to high missing values")
else:
    print(f"✅ No columns dropped due to high missing values based on criteria.")

print(f"DataFrame shape after dropping columns: {df_cleaned.shape}")

# Impute missing numerical values
print("Imputing missing numerical values (median)...")
for col in NUMERICAL_IMPUTE_COLUMNS:
    if col in df_cleaned.columns and df_cleaned[col].isnull().any():
        if pd.api.types.is_numeric_dtype(df_cleaned[col]):
            median_val = df_cleaned[col].median()
            df_cleaned[col] = df_cleaned[col].fillna(median_val)
            print(f" - Imputed '{col}' with median: {median_val}")
        else:
            print(f" - Warning: '{col}' is not numeric, skipping median imputation")
    elif col in df_cleaned.columns:
        print(f" - '{col}' has no missing values, skipping imputation.")
    else:
        print(f" - Warning: Numerical imputation column '{col}' not found in DataFrame")

# Impute missing categorical values
print("Imputing missing categorical values (mode)...")
for col in CATEGORICAL_IMPUTE_COLUMNS:
    if col in df_cleaned.columns and df_cleaned[col].isnull().any():
        # Check if mode is empty
        if not df_cleaned[col].mode().empty:
            mode_val = df_cleaned[col].mode()[0]
            df_cleaned[col] = df_cleaned[col].fillna(mode_val)
            print(f" - Imputed '{col}' with mode: '{mode_val}'")
        else:
            print(f" - Warning: Cannot find mode for '{col}', column might be entirely missing")
    elif col in df_cleaned.columns:
        print(f" - '{col}' has no missing values, skipping imputation.")
    else:
        print(f" - Warning: Categorical imputation column '{col}' not found in DataFrame")

print(f"✅ Data cleaning complete.")
print(f"DataFrame shape after cleaning: {df_cleaned.shape}")
print(f"Missing values after cleaning:\n", df_cleaned.isnull().sum())
print(f"- * 50)

# 4) Features & Target Preparation

print("\nPreparing features and target...")

# Verify all MODEL_FEATURES exist in the cleaned DataFrame
missing_model_features = [col for col in MODEL_FEATURES if col not in df_cleaned.columns]
if missing_model_features:
    print(f"❌ Error: The following model features are missing from the cleaned DataFrame: {missing_model_features}")
    exit()

if TARGET_COLUMN not in df_cleaned.columns:
    print(f"❌ Error: Target column '{TARGET_COLUMN}' not found in the DataFrame.")

```

```

exit()

X = df_cleaned[MODEL_FEATURES].copy()
y = df_cleaned[TARGET_COLUMN].copy()

# Handle potential NaN values in the target variable
if y.isnull().any():
    print(f"Warning: Target column '{TARGET_COLUMN}' contains missing values. Rows
    # Drop rows where the target is NaN
    initial_rows = X.shape[0]
    valid_indices = y.dropna().index
    X = X.loc[valid_indices]
    y = y.loc[valid_indices]
    print(f"Dropped {initial_rows - X.shape[0]} rows due to missing target values.")

print(f"Features (X) shape: {X.shape}")
print(f"Target (y) shape: {y.shape}")

# categorical features are string type for OneHotEncoder
print("Ensuring categorical features are string type...")
# Dynamically determine which of MODEL_FEATURES are categorical based on their incl
# cat_features for the ColumnTransformer are derived from MODEL_FEATURES that are m
cat_features_for_ohe = [col for col in MODEL_FEATURES if col in CATEGORICAL_IMPUTE_

for col in cat_features_for_ohe:
    if col in X.columns:
        X[col] = X[col].astype(str)
        print(f" - Converted '{col}' to string type.")
    else:
        print(f" - Warning: Categorical feature '{col}' not found in X, skipping t

# Identify numerical features for the preprocessor
num_features_for_scaler = [col for col in MODEL_FEATURES if col not in cat_features
# A more robust way:
# num_features_for_scaler = X.select_dtypes(include=np.number).columns.tolist()

print("✅ Features and target prepared.")
print("-" * 50)

# 5) Train-Test Split

print("\nSplitting data into training and testing sets...")
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=TEST_SIZE, random_state=RANDOM_STATE
)
print(f"✅ Data split complete: Training ({len(X_train)} samples), Testing ({len(X_
print("-" * 50)

# 6) Preprocessing (ColumnTransformer)

print("\nSetting up preprocessing pipeline (ColumnTransformer)...")
preprocessor = ColumnTransformer(transformers=[
    ('num', StandardScaler(), num_features_for_scaler),

```

```

    ('cat', OneHotEncoder(handle_unknown='ignore', sparse_output=False), cat_features)
])
print("✅ Preprocessor configured.")
print("-" * 50)

# 7) Model Pipeline

print("\nBuilding the full model pipeline...")
model = Pipeline(steps=[
    ('preprocessor', preprocessor),
    ('regressor', RandomForestRegressor(n_estimators=N_ESTIMATORS, random_state=RAN
)])
print("✅ Model pipeline created.")
print("-" * 50)

# 8) Train Model

print("\nStarting model training...")
try:
    model.fit(X_train, y_train)
    print("✅ Model training complete!")
except Exception as e:
    print(f"❌ An error occurred during model training: {e}")
    exit()
print("-" * 50)

# 9) Evaluate Model

print("\n🔄 Evaluating Model Performance...")
try:
    y_pred = model.predict(X_test)

    mae = mean_absolute_error(y_test, y_pred)

    # Calculate Mean Squared Error first
    mse = mean_squared_error(y_test, y_pred)
    # Then calculate Root Mean Squared Error by taking the square root of MSE
    rmse = np.sqrt(mse)

    r2 = r2_score(y_test, y_pred)

    print("\n✅ Model Evaluation Results:")
    print(f"    MAE: {mae:,.2f}")
    print(f"    RMSE: {rmse:,.2f}")
    print(f"    R²:   {r2:,.4f}")
    print("-" * 50)

except Exception as e:
    print(f"❌ An error occurred during model evaluation: {e}")
    exit()

print("\nScript execution finished.")

```

Attempting to load data...

✅ Data loaded successfully!

Initial DataFrame shape: (110490, 22)

Initial DataFrame head:

	TRANSACTION_NUMBER	INSTANCE_DATE	GROUP_EN	PROCEDURE_EN	\
0	102-1-2025	2025-01-07 15:57:40	Sales	Sell - Pre registration	
1	102-10-2025	2025-01-02 08:23:00	Sales	Sell - Pre registration	
2	102-100-2025	2025-01-02 15:18:51	Sales	Sell - Pre registration	
3	102-1000-2025	2025-01-08 11:02:57	Sales	Sell - Pre registration	
4	102-10000-2025	2025-02-07 11:15:01	Sales	Sell - Pre registration	

	IS_OFFPLAN_EN	IS_FREE_HOLD_EN	USAGE_EN	AREA_EN	PROP_TYPE_EN	\
0	Off-Plan	Free Hold	Residential	Wadi Al Safa 4	Unit	
1	Off-Plan	Free Hold	Residential	DUBAI SCIENCE PARK	Unit	
2	Off-Plan	Free Hold	Residential	AL FURJAN	Unit	
3	Off-Plan	Free Hold	Residential	Wadi Al Safa 5	Unit	
4	Off-Plan	Free Hold	Residential	Al Yelayiss 2	Unit	

	PROP_SB_TYPE_EN	...	ACTUAL_AREA	ROOMS_EN	PARKING	\
0	Flat	...	68.67	1 B/R	1	
1	Flat	...	47.25	Studio	1	
2	Hotel Apartment	...	32.35	Studio	1	
3	Flat	...	259.03	2 B/R	2	
4	Flat	...	66.08	1 B/R	1	

	NEAREST_METRO_EN	NEAREST_MALL_EN	\
0	NaN	NaN	
1	First Abu Dhabi Bank Metro Station	Mall of the Emirates	
2	ENERGY Metro Station	Ibn-e-Battuta Mall	
3	NaN	NaN	
4	NaN	NaN	

	NEAREST_LANDMARK_EN	TOTAL_BUYER	TOTAL_SELLER	MASTER_PROJECT_EN	\
0	NaN	0	0	NaN	
1	Motor City	0	0	NaN	
2	Expo 2020 Site	0	0	NaN	
3	NaN	0	0	NaN	
4	NaN	0	0	NaN	

	PROJECT_EN
0	Lacina
1	Binghatti Hills
2	AZIZI JEWEL
3	Verdes by Haven 1
4	The Baltimore

[5 rows x 22 columns]

Initial DataFrame info:

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 110490 entries, 0 to 110489

Data columns (total 22 columns):

#	Column	Non-Null Count	Dtype
0	TRANSACTION_NUMBER	110490 non-null	object
1	INSTANCE_DATE	110490 non-null	datetime64[ns]
2	GROUP_EN	110490 non-null	object

```

3  PROCEDURE_EN      110490 non-null object
4  IS_OFFPLAN_EN     110490 non-null object
5  IS_FREE_HOLD_EN   110490 non-null object
6  USAGE_EN          110490 non-null object
7  AREA_EN           110490 non-null object
8  PROP_TYPE_EN      110490 non-null object
9  PROP_SB_TYPE_EN   104394 non-null object
10 TRANS_VALUE       110490 non-null float64
11 PROCEDURE_AREA    110489 non-null float64
12 ACTUAL_AREA       110490 non-null float64
13 ROOMS_EN          94277 non-null object
14 PARKING            85646 non-null object
15 NEAREST_METRO_EN  56461 non-null object
16 NEAREST_MALL_EN   55007 non-null object
17 NEAREST_LANDMARK_EN 68028 non-null object
18 TOTAL_BUYER       110490 non-null int64
19 TOTAL_SELLER      110490 non-null int64
20 MASTER_PROJECT_EN 277 non-null object
21 PROJECT_EN        100900 non-null object
dtypes: datetime64[ns](1), float64(3), int64(2), object(16)
memory usage: 18.5+ MB
-----

```

Starting data cleaning...

Checking columns to drop with > 30% missing values...

- Dropping 'MASTER_PROJECT_EN' (missing: 99.75%)
- Dropping 'NEAREST_METRO_EN' (missing: 48.90%)
- Dropping 'NEAREST_MALL_EN' (missing: 50.22%)
- Dropping 'NEAREST_LANDMARK_EN' (missing: 38.43%)

✅ Dropped 4 columns due to high missing values.

DataFrame shape after dropping columns: (110490, 18)

Imputing missing numerical values (median)...

- Imputed 'PROCEDURE_AREA' with median: 86.6

Imputing missing categorical values (mode)...

- Imputed 'PARKING' with mode: '1'
- Imputed 'ROOMS_EN' with mode: '1 B/R'
- Imputed 'PROJECT_EN' with mode: 'Binghatti Skyrise'
- Imputed 'PROP_SB_TYPE_EN' with mode: 'Flat'

✅ Data cleaning complete.

DataFrame shape after cleaning: (110490, 18)

Missing values after cleaning:

```

TRANSACTION_NUMBER    0
INSTANCE_DATE          0
GROUP_EN               0
PROCEDURE_EN           0
IS_OFFPLAN_EN          0
IS_FREE_HOLD_EN        0
USAGE_EN               0
AREA_EN                0
PROP_TYPE_EN           0
PROP_SB_TYPE_EN        0
TRANS_VALUE            0
PROCEDURE_AREA         0
ACTUAL_AREA            0
ROOMS_EN               0
PARKING                0

```

```
TOTAL_BUYER      0
TOTAL_SELLER     0
PROJECT_EN       0
dtype: int64
```

Preparing features and target...

Features (X) shape: (110490, 5)

Target (y) shape: (110490,)

Ensuring categorical features are string type...

- Converted 'PARKING' to string type.
- Converted 'ROOMS_EN' to string type.
- Converted 'PROJECT_EN' to string type.
- Converted 'PROP_SB_TYPE_EN' to string type.

✅ Features and target prepared.

Splitting data into training and testing sets...

✅ Data split complete: Training (88392 samples), Testing (22098 samples).

Setting up preprocessing pipeline (ColumnTransformer)...

✅ Preprocessor configured.

Building the full model pipeline...

✅ Model pipeline created.

Starting model training...

✅ Model training complete!

🔄 Evaluating Model Performance...

✅ Model Evaluation Results:

MAE: 493,806.29

RMSE: 6,586,172.42

R²: 0.6660

Script execution finished.

In []: